

Assignment_2

January 28, 2024

1 Problem 1: Linear Regression from Scratch (30 points)

```
[1]: # import the necessary packages
import numpy as np
from matplotlib import pyplot as plt
np.random.seed(100)
```

Let's generate some data points first, by the equation $y = x - 3$.


```
[2]: x = np.random.randint(100, size=100)/30 - 2
X = x.reshape(-1, 1)

y = x + -3 + 0.3*np.random.randn(100)
```

Let's then visualize the data points we just created.

```
[3]: plt.scatter(X, y)
```

```
[3]: <matplotlib.collections.PathCollection at 0x7fb8930e2430>
```



1.1 1.1 Gradient of vanilla linear regression model (5 points)

In the lecture, we learn that the cost function of a linear regression model can be expressed as **Equation 1**:

$$J(\theta) = \frac{1}{2m} \sum_i^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

The gradient of it can be written as **Equation 2**:

$$\frac{\partial J(\theta)}{\partial \theta} = \frac{1}{m} \sum_i^m (h_{\theta}(x^{(i)}) - y^{(i)}) (x^{(i)})$$

1.2 1.2 Gradient of vanilla regularized regression model (5 points)

After adding the L2 regularization term, the linear regression model can be expressed as **Equation 3**:

$$J(\theta) = \frac{1}{2m} \sum_i^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m} \sum_j^n (\theta_j)^2$$

The gradient of it can be written as **Equation 4**:

$$\frac{\partial J(\theta)}{\partial \theta} = \frac{1}{m} \left(\sum_i^m (h_{\theta}(x^{(i)}) - y^{(i)}) (x^{(i)}) + \lambda \sum_j^n \theta_j \right)$$

1.3 Implement the cost function of a regularized regression model (5 points)

Please implement the cost function of a regularized regression model according to the above equations.

1.4 Implement the gradient of the cost function of a regularized regression model (5 points)

Please implement the gradient of the cost function of a regularized regression model according to the above equations.

```
[4]: def regularized_linear_regression(X, y, alpha=0.01, lambda_value=1, epochs=30):
    """
    :param x: feature matrix
    :param y: target vector
    :param alpha: learning rate (default:0.01)
    :param lambda_value: lambda (default:1)
    :param epochs: maximum number of iterations of the
                    linear regression algorithm for a single run (default=30)
    :return: weights, list of the cost function changing overtime
    """

    m = np.shape(X)[0] # total number of samples
    n = np.shape(X)[1] # total number of features

    X = np.concatenate((np.ones((m, 1)), X), axis=1)
    W = np.random.randn(n + 1, )

    # stores the updates on the cost function (loss function)
    cost_history_list = []

    # iterate until the maximum number of epochs
    for current_iteration in np.arange(epochs): # begin the process

        # compute the dot product between our feature 'X' and weight 'W'
        y_estimated = X.dot(W)

        # calculate the difference between the actual and predicted value
        error = y_estimated - y
```

#####

↪

```
##### Begin of Question 1.3
```

```
→#####
```

```
#####
```

```
##### Please write down your code here:####
```

```
# calculate the cost (MSE) (Equation 1)
```

```
cost_without_regularization = (1/(2*m))*np.sum(np.power(error, 2))
```

```
##### Please write down your code here:####
```

```
# regularization term
```

```
reg_term = (lambda_value/(2*m))*np.sum(np.square(W))
```

```
# calculate the cost (MSE) + regularization term (Equation 3)
```

```
cost_with_regularization = cost_without_regularization + reg_term
```

```
#####
```

```
→
```

```
##### End of Question 1.3
```

```
→#####
```

```
#####
```

```
→
```

```
#####
```

```
→
```

```
##### Begin of Question 1.4
```

```
→#####
```

```
#####
```

```
→
```

```
##### Please write down your code here:####
```

```
# calculate the gradient of the cost function with regularization term
```

```
→(Equation )
```

```
gradient = (1/m)*(X.T.dot(error)+(lambda_value*W))
```

```
# Now we have to update our weights
```

```
W = W - alpha * gradient
```

```
#####
```

```
→
```

```
##### End of Question 1.4
#####
#####

# keep track the cost as it changes in each iteration
cost_history_list.append(cost_with_regularization)

# Let's print out the cost
print(f"Cost with regularization: {cost_with_regularization}")
print(f"Mean square error: {cost_without_regularization}")

return W, cost_history_list
```

Run the following code to train your model.


Hint: If you have the correct code written above, the cost should be 0.5181222986588751 when $\lambda = 10$.

```
[5]: lambda_list = [0, 10, 100, 1000, 10000]
for lambda_ in lambda_list:
    # calls regression function with different values of lambda
    weight, _ = regularized_linear_regression(X, y, alpha=0.01,
                                              lambda_value=lambda_, epochs=1000)


    fitted_line = np.dot(X, weight[1]) + weight[0]
    plt.scatter(X, y, label='data points')
    plt.plot(X, fitted_line, color='r', label='Fitted line')
    plt.xlabel("X")
    plt.ylabel("y")
    plt.title(f"Regression (lambda : {lambda_})")
    plt.legend()
    plt.show()
```

Cost with regularization: 0.05165888565058274


Mean square error: 0.05165888565058274




Cost with regularization: 0.5181225049184746
Mean square error: 0.08982014821513126




Cost with regularization: 2.793172488740026
Mean square error: 1.2785107029715972



Cost with regularization: 5.591464362606628
Mean square error: 4.946888025066496



Cost with regularization: 6.2426956269339735
Mean square error: 6.1614425833558135



1.5 Analyze your results (10 points)

According to the above figures, what's the best choice of λ ?

Why the regressed line turns to be flat as we increase λ ?

Your answer: 1. $\lambda = 0$

2. In regularized linear regression, increasing the value of the lambda tends to shrink the regression coefficient to zero. When the lambda value increases to a really large value, the magnitude of the regression coefficients should be very small, resulting in a flat regression line that is close to parallel to the x-axis.

2 Problem 2: Getting familiar with PyTorch (30 points)

```
[6]: import mltools as ml
import torch

data = np.genfromtxt("data/curve80.txt")
X = data[:,0]
```

```

X = np.atleast_2d(X).T # code expects shape (M,N) so make sure it's ↵
    ↪2-dimensional
Y = data[:,1] # doesn't matter for Y
Xtr,Xte,Ytr,Yte = ml.splitData(X,Y,0.75) # split data set 75/25
degree = 5
XtrP = ml.transforms.fpoly(Xtr, degree=degree, bias=False)
XtrP, params = ml.transforms.rescale(XtrP)

XtrP_tensor=torch.from_numpy(XtrP).float()
Ytr_tensor=torch.from_numpy(Ytr).reshape(-1,1).float()
XtrP_tensor=XtrP_tensor.view(-1,degree)
Ytr_tensor=Ytr_tensor.view(-1,1)

```

```

[7]: linear_regressor = torch.nn.Linear(in_features=degree, out_features=1)
    print(linear_regressor)

```

```

Linear(in_features=5, out_features=1, bias=True)

```

```

[8]: criterion = torch.nn.MSELoss()
    optimizer = torch.optim.SGD(linear_regressor.parameters(), lr=0.1)
    epochs = 100000

```

```

[9]: loss_record = []
    for _ in range(epochs):
        optimizer.zero_grad() # set gradient to zero
        pred_y = linear_regressor(XtrP_tensor)
        loss = criterion(pred_y, Ytr_tensor) # calculate loss function
        loss.backward() # backpropagate gradient
        loss_record.append(loss.item())
        optimizer.step() # update the parameters in the linear ↵
    ↪regressor

```

```


[10]: plt.plot(range(epochs), (loss_record))

```

```


[10]: [<matplotlib.lines.Line2D at 0x7fb89b0c1ca0>]

```



```
[11]: xs = np.linspace(0,10,200)
xs = xs[:,np.newaxis]
xsP, _ = ml.transforms.rescale(ml.transforms.
    ↳fpoly(xs,degree=degree,bias=False), params)
xsP_tensor = torch.from_numpy(xsP).float()
ys = linear_regressor(xsP_tensor)
plt.scatter(Xtr,Ytr,label="training data")
plt.plot(xs,ys.detach().numpy(),label="prediction function",color='red')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
```

```
[11]: <matplotlib.legend.Legend at 0x7fb89b46d9a0>
```



3 Statement of Collaboration

I do it by myself