

Assignment_1

January 21, 2024

1 Problem 1: Python & Data Exploration



```
[2]: import numpy as np
import matplotlib.pyplot as plt
iris = np.genfromtxt("data/iris.txt", delimiter=None) # load the text file
Y = iris[:, -1] # target value is the last column
X = iris[:, 0:-1] # features are the other columns


#1.1
num_features = X.shape[1]
num_dataPoints = X.shape[0]
print("Problem1.1:")
print("Number of features: {0}".format(num_features))
print("Number of data points: {0}".format(num_dataPoints))
```


Problem1.1:
Number of features: 4
Number of data points: 148

```
[2]: # 1.2
print("Problem1.2:")
for i in range(0, num_features):
    plt.hist(X[:, i], label='feature{}'.format(i+1))
    plt.legend()
    plt.show()
    plt.close()
```

Problem1.2:







```
[4]: #1.3
print("Problem1.3:")
means = np.mean(X, axis=0)
std_devs = np.std(X, axis=0)
for i in range(0, num_features):
    print("feature{:}: mean: {} standard: {}".format(i, means[i], std_devs[i]))
```

Problem1.3:

```
feature0: mean: 5.900103764189187 standard: 0.8334020667748939
feature1: mean: 3.0989309168918915 standard: 0.43629183800107685
feature2: mean: 3.8195548405405413 standard: 1.7540571093439352
feature3: mean: 1.252555484594594 standard: 0.7587724570263246
```

```
[4]: #1.4
print("Problem1.4:")
colors=[]
for y in Y:
    if y==0:
        colors.append('red')
    elif y==1:
        colors.append('blue')
    else:
```


```


        colors.append('green')


for i in range(1, num_features):
    f1=X[:,0]
    f2=X[:,i]
    plt.scatter(f1,f2,c=colors)
    plt.xlabel('feature1')
    plt.ylabel('feature{}'.format(i+1))
    plt.title('Scatter plot of Feature1 vs Feature{}'.format(i+1))
    plt.show()

```

Problem1.4:







2 Problem 2: k-Nearest Neighbor (kNN) exercise

The kNN classifier consists of two stages:

- During training, the classifier takes the training data and simply remembers it
- During testing, kNN classifies every test image by comparing to all training images and transferring the labels of the k most similar training examples
- The value of k is cross-validated

In this exercise you will implement these steps and understand the basic Image Classification pipeline, cross-validation, and gain proficiency in writing efficient, vectorized code.

```
[5]: # Run some setup code for this notebook.  
  
import random  
import numpy as np  
from cs273p.data_utils import load_CIFAR10  
import matplotlib.pyplot as plt
```

```
# This is a bit of magic to make matplotlib figures appear inline in the
↳notebook
# rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/
↳autoreload-of-modules-in-ipython
# %load_ext autoreload
# %autoreload 2
```

```
[6]: %cd cs273p/datasets
!source get_datasets.sh
```

```
/Users/huangjiayi/Documents/003UCI/learn/273/hw/hw1/cs273p/datasets
/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-
packages/IPython/core/magics/osm.py:417: UserWarning: using dhist requires you
to install the `pickleshare` library.
  self.shell.db['dhist'] = compress_dhist(dhist)[-100:]
--2024-01-20 14:59:37-- http://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
  www.cs.toronto.edu (www.cs.toronto.edu)... 128.100.3.30
  www.cs.toronto.edu (www.cs.toronto.edu)|128.100.3.30|:80...
  HTTP ... 200 OK
  170498071 (163M) [application/x-gzip]
  : "cifar-10-python.tar.gz"

cifar-10-python.tar 100%[=====>] 162.60M  9.15MB/s    21s

2024-01-20 14:59:58 (7.75 MB/s) - "cifar-10-python.tar.gz"
[170498071/170498071])

x cifar-10-batches-py/
x cifar-10-batches-py/data_batch_4
x cifar-10-batches-py/readme.html
x cifar-10-batches-py/test_batch
x cifar-10-batches-py/data_batch_3
x cifar-10-batches-py/batches.meta
x cifar-10-batches-py/data_batch_2
x cifar-10-batches-py/data_batch_5
x cifar-10-batches-py/data_batch_1
```

```
[7]: %cd ../..
```

```
/Users/huangjiayi/Documents/003UCI/learn/273/hw/hw1
```

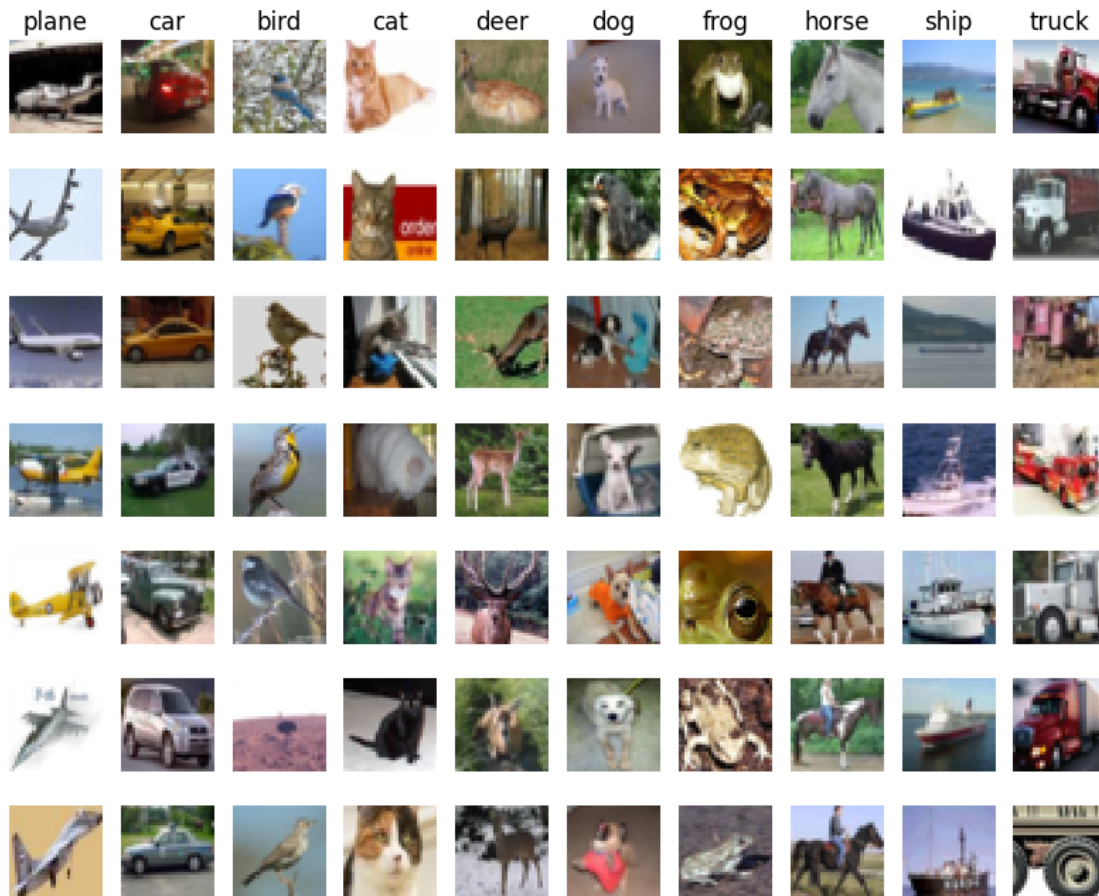


```
[8]: # Load the raw CIFAR-10 data.
cifar10_dir = './cs273p/datasets/cifar-10-batches-py'
X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

```
[9]: # Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', '
↳ ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
```



```
[10]: # Subsample the data for more efficient code execution in this exercise
num_training = 5000
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]

num_test = 500
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]
```

```
[11]: # Reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
print(X_train.shape, X_test.shape)
```

```
(5000, 3072) (500, 3072)
```

```
[12]: from cs273p.classifiers import KNearestNeighbor

# Create a kNN classifier instance.
# Remember that training a kNN classifier is a noop:
# the Classifier simply remembers the data and does no further processing
classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
```

We would now like to classify the test data with the kNN classifier. Recall that we can break down this process into two steps:

1. First we must compute the distances between all test examples and all train examples.
2. Given these distances, for each test example we find the k nearest examples and have them vote for the label

Lets begin with computing the distance matrix between all training and test examples. For example, if there are N_{tr} training examples and N_{te} test examples, this stage should result in a $N_{te} \times N_{tr}$ matrix where each element (i,j) is the distance between the i-th test and j-th train example.


First, open `cs273p/classifiers/k_nearest_neighbor.py` and implement the function `compute_distances_two_loops` that uses a (very inefficient) double loop over all pairs of (test, train) examples and computes the distance matrix one element at a time.

```
[13]: # Open cs273p/classifiers/k_nearest_neighbor.py and implement
# compute_distances_two_loops.

# Test your implementation:
dists = classifier.compute_distances_two_loops(X_test)
print(dists.shape)
```

(500, 5000)

```
[14]: # We can visualize the distance matrix: each row is a single test example and
# its distances to training examples
plt.imshow(dists, interpolation='none')
plt.show()
```



Inline Question #1: Notice the structured patterns in the distance matrix, where some rows or columns are visible brighter. (Note that with the default color scheme black indicates low distances while white indicates high distances.)

- What in the data is the cause behind the distinctly bright rows?
- What causes the columns?

Your Answer: 1. There are outliers in the test data compared to the training data 2. There are outliers in the training data compared to the test data

```
[15]: # Now implement the function predict_labels and run the code below:
# We use k = 1 (which is Nearest Neighbor).
y_test_pred = classifier.predict_labels(dists, k=1)

# Compute and print the fraction of correctly predicted examples
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 137 / 500 correct => accuracy: 0.274000

You should expect to see approximately 27% accuracy. Now lets try out a larger k, say k = 5:

```
[16]: y_test_pred = classifier.predict_labels(dists, k=5)
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 139 / 500 correct => accuracy: 0.278000

You should expect to see a slightly better performance than with k = 1.

```
[17]: # Now lets speed up distance matrix computation by using partial vectorization
# with one loop. Implement the function compute_distances_one_loop and run the
# code below:
dists_one = classifier.compute_distances_one_loop(X_test)

# To ensure that our vectorized implementation is correct, we make sure that it
# agrees with the naive implementation. There are many ways to decide whether
# two matrices are similar; one of the simplest is the Frobenius norm. In case
# you haven't seen it before, the Frobenius norm of two matrices is the square
# root of the squared sum of differences of all elements; in other words,
# ↪ reshape
# the matrices into vectors and compute the Euclidean distance between them.
difference = np.linalg.norm(dists - dists_one, ord='fro')
print('Difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')
```

Difference was: 0.000000

Good! The distance matrices are the same

```
[18]: # Now implement the fully vectorized version inside compute_distances_no_loops
# and run the code
dists_two = classifier.compute_distances_no_loops(X_test)
```

```

# check that the distance matrix agrees with the one we computed before:
difference = np.linalg.norm(dists - dists_two, ord='fro')
print('Difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')

```

Difference was: 0.000000

Good! The distance matrices are the same

```

[19]: # Let's compare how fast the implementations are
def time_function(f, *args):
    """
    Call a function f with args and return the time (in seconds) that it took
    to execute.
    """
    import time
    tic = time.time()
    f(*args)
    toc = time.time()
    return toc - tic

two_loop_time = time_function(classifier.compute_distances_two_loops, X_test)
print('Two loop version took %f seconds' % two_loop_time)

one_loop_time = time_function(classifier.compute_distances_one_loop, X_test)
print('One loop version took %f seconds' % one_loop_time)

no_loop_time = time_function(classifier.compute_distances_no_loops, X_test)
print('No loop version took %f seconds' % no_loop_time)

# you should see significantly faster performance with the fully vectorized
implementation

```

Two loop version took 92.284572 seconds

One loop version took 68.989297 seconds

No loop version took 1.032221 seconds

2.0.1 Cross-validation

We have implemented the k-Nearest Neighbor classifier but we set the value $k = 5$ arbitrarily. We will now determine the best value of this hyperparameter with cross-validation.

```

[20]: num_folds = 5
      k_choices = [1, 3, 5, 8, 10, 12, 15, 20, 50, 100]

```

```

X_train_folds = []
y_train_folds = []
#####
# TODO:
# Split up the training data into folds. After splitting, X_train_folds and
# y_train_folds should each be lists of length num_folds, where
# y_train_folds[i] is the label vector for the points in X_train_folds[i].
# Hint: Look up the numpy array_split function.
#####
X_train_folds = np.array_split(X_train, num_folds)
y_train_folds = np.array_split(y_train, num_folds)
#####
#                                     END OF YOUR CODE
#####

# A dictionary holding the accuracies for different values of k that we find
# when running cross-validation. After running cross-validation,
# k_to_accuracies[k] should be a list of length num_folds giving the different
# accuracy values that we found when using that value of k.
k_to_accuracies = {}

#####
# TODO:
# Perform k-fold cross validation to find the best value of k. For each
# possible value of k, run the k-nearest-neighbor algorithm num_folds times,
# where in each case you use all but one of the folds as training data and the
# last fold as a validation set. Store the accuracies for all fold and all
# values of k in the k_to_accuracies dictionary.
#####

for i in range(num_folds):
    # Create training data by combining all folds except the current one
    X_training = np.concatenate([fold for j, fold in enumerate(X_train_folds)
    ↪if j != i])
    y_training = np.concatenate([fold for j, fold in enumerate(y_train_folds)
    ↪if j != i])

    # Create validation data using the current fold
    X_val_cv = X_train_folds[i]
    y_val_cv = y_train_folds[i]

    # Initialize and train the k-nearest-neighbor classifier
    knn_classifier = KNearestNeighbor()
    knn_classifier.train(X_training, y_training)
    dists = knn_classifier.compute_distances_no_loops(X_val_cv)

```

```

for k in k_choices:
    if i == 0:
        k_to_accuracies[k] = []

        # Compute accuracy on the validation set and store in k_to_accuracies
        predictions= knn_classifier.predict_labels(dists, k=k)
        correct_predictions = np.sum(predictions == y_val_cv)
        accuracy = correct_predictions / len(y_val_cv)
        k_to_accuracies[k].append(accuracy)
#####
#                                     END OF YOUR CODE                                     #
#####

# Print out the computed accuracies
for k in sorted(k_to_accuracies):
    for accuracy in k_to_accuracies[k]:
        print('k = %d, accuracy = %f' % (k, accuracy))

```

```

k = 1, accuracy = 0.263000
k = 1, accuracy = 0.257000
k = 1, accuracy = 0.264000
k = 1, accuracy = 0.278000
k = 1, accuracy = 0.266000
k = 3, accuracy = 0.239000
k = 3, accuracy = 0.249000
k = 3, accuracy = 0.240000
k = 3, accuracy = 0.266000
k = 3, accuracy = 0.254000
k = 5, accuracy = 0.248000
k = 5, accuracy = 0.266000
k = 5, accuracy = 0.280000
k = 5, accuracy = 0.292000
k = 5, accuracy = 0.280000
k = 8, accuracy = 0.262000
k = 8, accuracy = 0.282000
k = 8, accuracy = 0.273000
k = 8, accuracy = 0.290000
k = 8, accuracy = 0.273000
k = 10, accuracy = 0.265000
k = 10, accuracy = 0.296000
k = 10, accuracy = 0.276000
k = 10, accuracy = 0.284000
k = 10, accuracy = 0.280000
k = 12, accuracy = 0.260000
k = 12, accuracy = 0.295000
k = 12, accuracy = 0.279000
k = 12, accuracy = 0.283000
k = 12, accuracy = 0.280000

```

```

k = 15, accuracy = 0.252000
k = 15, accuracy = 0.289000
k = 15, accuracy = 0.278000
k = 15, accuracy = 0.282000
k = 15, accuracy = 0.274000
k = 20, accuracy = 0.270000
k = 20, accuracy = 0.279000
k = 20, accuracy = 0.279000
k = 20, accuracy = 0.282000
k = 20, accuracy = 0.285000
k = 50, accuracy = 0.271000
k = 50, accuracy = 0.288000
k = 50, accuracy = 0.278000
k = 50, accuracy = 0.269000
k = 50, accuracy = 0.266000
k = 100, accuracy = 0.256000
k = 100, accuracy = 0.270000
k = 100, accuracy = 0.263000
k = 100, accuracy = 0.256000
k = 100, accuracy = 0.263000


```

```

[21]: # plot the raw observations
for k in k_choices:
    accuracies = k_to_accuracies[k]
    plt.scatter([k] * len(accuracies), accuracies)

# plot the trend line with error bars that correspond to standard deviation
accuracies_mean = np.array([np.mean(v) for k,v in sorted(k_to_accuracies.
    ↪items())])
accuracies_std = np.array([np.std(v) for k,v in sorted(k_to_accuracies.
    ↪items())])
plt.errorbar(k_choices, accuracies_mean, yerr=accuracies_std)
plt.title('Cross-validation on k')
plt.xlabel('k')
plt.ylabel('Cross-validation accuracy')
plt.show()

```

```
[22]: # Based on the cross-validation results above, choose the best value for k,
# retrain the classifier using all the training data, and test it on the test
# data. You should be able to get above 28% accuracy on the test data.
best_k = 10

classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
y_test_pred = classifier.predict(X_test, k=best_k)

# Compute and display the accuracy
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))

Got 141 / 500 correct => accuracy: 0.282000
```

3 Problem 3: Naïve Bayes Classifiers

3.1

$$P(y=-1) = 3/5$$

$$P(y=1) = 2/5$$

$$P(x_1=0 \mid y=-1) = 1/2$$

$$P(x_1=0 \mid y=1) = 1/4$$

$$P(x_1=1 \mid y=-1) = 1/2$$

$$P(x_1=1 \mid y=1) = 3/4$$

$$P(x_2=0 \mid y=-1) = 1/6$$

$$P(x_2=0 \mid y=1) = 1$$

$$P(x_2=1 \mid y=-1) = 5/6$$

$$P(x_2=1 \mid y=1) = 0$$

$$P(x_3=0 \mid y=-1) = 1/3$$

$$P(x_3=0 \mid y=1) = 1/4$$

$$P(x_3=1 \mid y=-1) = 2/3$$

$$P(x_3=1 \mid y=1) = 3/4$$

$$P(x_4=0 \mid y=-1) = 1/6$$

$$P(x_4=0 \mid y=1) = 1/2$$

$$P(x_4=1 \mid y=-1) = 5/6$$

$$P(x_4=1 \mid y=1) = 1/2$$

$$P(x_5=0 \mid y=-1) = 2/3$$

$$P(x_5=0 \mid y=1) = 3/4$$

$$P(x_5=1 \mid y=-1) = 1/3$$

$$P(x_5=1 \mid y=1) = 1/4$$

3.2

$$1.X = (0 \ 0 \ 0 \ 0 \ 0)$$

$$P(y|x) = P(x_1|y) * P(x_2|y) * P(x_3|y) * P(x_4|y) * P(x_5|y) * P(y)$$

For $y=1$:

$$P(y=1 \mid X = (0 \ 0 \ 0 \ 0 \ 0))$$

$$= (1/4) * 1 * (1/4) * (1/2) * (3/4) * (2/5)$$

$$= 0.009375$$

For $y=-1$:

$$P(y=-1 \mid X = (0 \ 0 \ 0 \ 0 \ 0))$$

$$= (1/2) * (1/6) * (1/3) * (1/6) * (2/3) * (3/5)$$

$$= 0.00185185$$

So the class predicted for $X = (0 \ 0 \ 0 \ 0 \ 0)$ is $y=1$

$$2.X = (1 \ 1 \ 0 \ 1 \ 0)$$

$$\text{For } y=1: P(y=1 \mid X = (1 \ 1 \ 0 \ 1 \ 0))$$

$$= (3/4) * 0 * 1/4 * (1/2) * (3/4) * (2/5) \\ = 0$$

For $y=-1$:

$$P(y=-1 | X = (1 \ 1 \ 0 \ 1 \ 0)) \\ = (1/2) * (5/6) * (1/3) * (5/6) * (2/3) * (3/5) \\ = 0.0462963$$

So the class predicted for $X = (1 \ 1 \ 0 \ 1 \ 0)$ is $y=-1$

3.3

$$P(y=1 | x=(1 \ 1 \ 0 \ 1 \ 1)) \\ = P(x=(1 \ 1 \ 0 \ 1 \ 1) | y=1) * P(y=1) / P(x=(1 \ 1 \ 0 \ 1 \ 0)) \\ = 0$$

3.4

The joint Bayes classifier involves estimating joint probabilities for all possible combinations of features. If we use a “joint” Bayes classifier, there will be $2^5=32$ probability to estimate classification probability.

3.5

Yes, we should re-train the model.

Assume the naïve Bayes model only have features $x_2 \ x_3 \ x_4 \ x_5$.

$$P(y | x_2 \ x_3 \ x_4 \ x_5) = P(y) * P(x_2|y) * P(x_3|y) * P(x_4|y) * P(x_5|y)$$

We can get the probability of $P(y) \ P(x_2|y) \ P(x_3|y) \ P(x_4|y) \ P(x_5|y)$.

4 Problem4: Statement of Collaboration

I do this work by myself.